Kaaliber aims to solve the problem of fragmented, expensive, and inefficient customer growth by bringing the entire marketing process into one AI-powered platform. Instead of managing multiple advertising channels, disconnected data, creative tools, and customer retention systems separately, Kaaliber unifies them into a single engine that helps businesses acquire high-value customers, automatically optimise budgets, continuously create fresh ads, understand what drives revenue, and retain existing users. By making real-time, data-driven decisions focused on actual business outcomes rather than just clicks and installs, Kaaliber aims to reduce wasted marketing spend, minimise reliance on multiple intermediaries, and help brands grow more efficiently and profitably across channels.

Key responsibilities

• Define AI/ML architecture, model-selection strategy, reusable service patterns and technical

standards across products.

• Lead the design of LLM workflows, retrieval systems, tool use, routing, structured outputs and

human-review controls.

• Guide predictive-model development for classification, ranking, recommendation, forecasting

and anomaly detection.

• Establish dataset, experimentation, offline evaluation, online measurement and model-release

practices.

• Define quality, safety, latency and cost thresholds and ensure model changes are validated

before production release.

• Design MLOps and LLMOps capabilities covering versioning, deployment, tracing, drift, rollback

and provider resilience.

• Review AI designs and code, create reference implementations and resolve complex modeling or

production issues.

• Mentor AI/ML engineers, improve technical depth and coordinate delivery with backend,

frontend, QA, security and cloud teams.

• Assess external models and platforms while protecting data, limiting vendor dependence and

controlling operational cost.

• Communicate evidence, uncertainty, model limitations and technical trade-offs clearly to

leadership and delivery teams.

Must-have capabilities

• Advanced Python and strong foundations in machine learning, deep learning, statistics and

experimental design alongwith 5+years of experience into ML domain.

• Proven ownership of production LLM applications, including structured outputs, tool calling,

retrieval and evaluation.

• Experience delivering classical ML systems using scikit-learn, PyTorch, TensorFlow or comparable

frameworks.

• Strong understanding of precision, recall, ranking metrics, calibration, drift, bias, data leakage

and error analysis.

• Experience with model deployment, versioning, feature or data pipelines, observability and

rollback.

• Ability to architect secure AI services and defend against prompt injection, unsafe tool use and

sensitive-data leakage.

• Strong SQL and data-analysis skills with experience working on event, behavioral or time-series

datasets.

• Technical leadership, mentoring, design review and stakeholder communication experience.

Preferred experience

• Multi-model routing, agent orchestration, vector retrieval and knowledge-graph experience.

• Causal inference, uplift modeling, recommendations, optimization or contextual bandits.

• Multimodal models involving text, image, audio or video.

• Cloud-based ML platforms, Kubernetes, distributed inference and GPU workload optimization.

• Experience building shared AI platforms or standards used by more than one product team.